

B16 Series B

Tumour scenarios: inferCNV and calling malignant cells


Separate malignant from normal cells with CNV signal — and know when the inference lies to you.

About 38 min | Prerequisites: B5-B10 | Data: infercnv built-in oligodendroglioma example

Three epithelial clusters — which is cancer?

simulated data

clusters: epithelium splits in three



T cells

macrophages

which cluster is malignant?

EPCAM: all three are positive



the marker says epithelial — not malignant

One tumour sample: three epithelial clusters, all EPCAM-positive. Clustering cannot help you

WHERE THIS EPISODE STARTS

**Malignancy is defined in the genome —
clustering and markers only see the transcriptome.**

To call malignancy you must dig the genome's shadow out of the RNA. That is exactly what inferCNV does.

What you will learn



- 1 Explain why clustering and markers cannot settle malignant vs normal in tumour data
- 2 Understand inferCNV's chain (genome ordering, sliding windows, reference-relative shifts) and run the built-in example
- 3 Read clone structure off the heatmap, triangulate a verdict, and name three common false signals

01

Where the signal comes from

scRNA never reads DNA — CNV here is inferred, not measured


The defining feature of malignancy is genomic



scRNA never reads DNA — but large gains and losses cast a shadow on segment-average expression

Large gains and losses cast a shadow on segment-average expression


The chain: order, window, baseline



same numbers: ordering lines the signal up, averaging crushes the noise (~1 log2 unit down to ~1/7)

The same numbers, three steps: genome ordering, windowed averaging, reference-relative shifts

The reference: who counts as normal



every colour on the heatmap is relative to this baseline — pick it wrong and the whole map is wrong

First choice: immune / stromal cells from the same sample — their genomes are almost surely normal

CORE IDEA

**The CNV signal is a statistical shadow:
single genes mean nothing — only segments count.**

One gene halving has a hundred explanations; a whole chromosome arm halving usually has one.

02

Three inputs and the example data

counts, annotations, gene order — all three required

inferCNV eats three files

1. counts matrix

	MGH36_A01	MGH53_B02
DLL3	0	12
EGFR	3	0
S0X4	8	21

raw gene-by-cell counts
(extractable from Seurat)

2. annotations

MGH36_A01	malignant_MGH36
MGH53_B02	malignant_MGH53
Mic_93_C07	Microglia/Macro.
Oli_97_D11	Oligodendrocytes

which group each cell belongs to
references are named groups

3. gene order file

WASH7P	chr1	14363
NOC2L	chr1	879584
SAMD11	chr1	860260
...	...	...

chromosome + position per gene
ordering and windows depend on it

`CreateInfercnvObject(counts, annotations, gene_order, ref_group_names)`

You already have counts; annotations come from your labels; the gene order file is the new one

Load the package and the built-in example

```
# BiocManager::install("infercnv") # HMM 功能需要先装 JAGS
library(infercnv)
set.seed(1234)

counts_f <- system.file("extdata",
  "oligodendroglioma_expression_downsampled_counts_matrix.gz",
  package = "infercnv")
anno_f <- system.file("extdata",
  "oligodendroglioma_annotations_downsampled.txt",
  package = "infercnv")
gene_f <- system.file("extdata",
  "encode_downsampled.EXAMPLE_ONLY_DONT_REUSE.txt",
  package = "infercnv")
```

system.file points at files shipped with the package — nothing to download

Create the infercnv object

```
anno <- read.table(anno_f, sep = "\t",
                  colnames = c("cell", "group"))
head(table(anno$group), 3)

cnv <- CreateInfercnvObject(
  raw_counts_matrix = counts_f,
  annotations_file = anno_f,
  gene_order_file = gene_f,
  delim = "\t",
  ref_group_names = c("Microglia/Macrophage",
                      "Oligodendrocytes (non-malignant)"))
```

Microglia/Macrophage	Oligodendrocytes (non-malignant)
27	22
malignant_MGH36	
33	

ref_group_names must match the annotation group names exactly — typos error out

What you have at this point

One matrix



Smart-seq
oligodendroglioma
counts: ~10k genes,
~180 cells

Two identities



two reference groups
(microglia,
oligodendrocytes) plus
tumour observation
groups

One map



chromosome and
position per gene —
ordering and windows
depend on it

03

Run it, and the parameters

cutoff, window and denoise each own one job


infercnv::run — first pass

```
cnv <- infercnv::run(cnv,  
  cutoff = 1,    # 平均計數低於它的基因剔除:  
                # Smart-seq 用 1, 10x 用 0.1  
  out_dir = "output/b16_infercnv",  
  cluster_by_groups = TRUE, # 觀察組各自聚類  
  denoise = TRUE,         # 壓掉基線附近的白噪音  
  HMM = FALSE)           # 第一輪先看連續訊號
```

```
INFO [...] ::process_data:Start  
INFO [...] STEP 15: computing tum or subclusters  
INFO [...] Making the final infercnv heatmap
```

It runs for minutes and logs hundreds of lines — the line that matters says the heatmap is done

Three parameters, three jobs



- cutoff drops low-expression genes — noise genes only dilute the window
- window (default 101 genes): too small keeps noise, too large blurs boundaries
- denoise flattens values within the reference's own wobble — a much cleaner map

After the run: inside out_dir

```
list.files("output/b16_infercnv") |> head(6)
```

```
[1] "infercnv.png"  
[2] "infercnv.preliminary.png"  
[3] "infercnv.references.txt"  
[4] "infercnv.observations.txt"  
[5] "infercnv.observation_groupings.txt"  
[6] "run.final.infercnv_obj"
```


infercnv.png is the main figure; run.final.infercnv_obj is the result object we will read back

04

Reading the heatmap

reference first, events second, clones last


Anatomy of the heatmap




The reference half should be near-white; in the observations, only contiguous blue/red blocks are events

The textbook event: chr1p/19q co-deletion

chr1: p arm drops, q arm returns to baseline



chr19: the q arm drops



reference



malignant mean



simulated data

1p/19q co-deletion is the textbook oligodendroglioma event — used clinically for typing

The molecularly defining event of oligodendroglioma — the real example data shows exactly these two

Subclones: the tumour's family tree

four chromosomes (subset of the real result)



shared events form the trunk; partial events form branches

simulated data

Shared events form the trunk, partial events the branches — row clustering lays it out for you

A FIXED READING ORDER

- 1. Is the reference clean?**
- 2. Any contiguous events?**
- 3. How do they stratify?**

If the reference is not clean, nothing downstream counts — the baseline itself is broken.

05

HMM mode and CopyKAT

from shades of colour to discrete calls


HMM mode: colours become states

```
cnv_hm m <- infercnv::run(cnv0, # 用未跑過的原始物件
  cutoff = 1,
  out_dir = "output/b16_infercnv_hm m ",
  cluster_by_groups = TRUE,
  denoise = TRUE,
  HMM = TRUE,
  HMM_type = "i6") # i3 = 缺失 / 中性 / 增加; i6 再分程度
```

HMM is several times slower — keep it off while exploring, turn it on to conclude

HMM states and per-cell calls

HMM: translate the continuous signal into discrete states (loss / neutral / gain)




called loss

called gain

subclone B mean signal + a real 3-state Viterbi

inferCNV's i3 mode = these three states; i6 splits gains further, turning shades of colour into integer copy-number calls

per-cell calls: two clean clouds



CopyKAT's idea: call aneuploid vs diploid per cell, no groups needed up front

Left: continuous signal translated into states. Right: CopyKAT's idea — call aneuploid vs diploid per cell

inferCNV vs CopyKAT

not rivals but complements — people often run both

inferCNV

group-level, great visuals

- you must name the reference groups
- the heatmap communicates; subclone structure is clear
- verdicts live at the group / cluster level

CopyKAT

per-cell, automatic baseline

- no reference needed — it finds diploid cells itself
- outputs aneuploid / diploid per cell
- shines when normal epithelium hides among tumour cells

06

The verdict: triangulate

CNV x clustering x markers, written into metadata

A per-cell CNV score from the result object

```
obj <- readRDS("output/b16_infercnv/run.final_infercnv_obj")
expr <- obj@expr.data      # 基因 × 細胞, 中性值在 1 附近
dim(expr)


cnv_score <- colMeans((expr - 1)^2) # 偏移平方的平均
ref_cells <- unlist(obj@reference_grouped_cell_indices)
summary(cnv_score[ref_cells])      # reference 的分數範圍
```

```
[1] 7113 184
```

Min.	1stQu.	Median	Mean	3rdQu.	Max.
0.000876	0.001392	0.001729	0.001846	0.002206	0.004215

expr.data is the smoothed, centred matrix — the score is each cell's distance from 1

The triangulation flow



write the verdict into metadata (malignant / normal / ambiguous) — downstream starts from that column

Any single line of evidence can lie; only the three-way cross-check gets the error rate acceptably low

Make the call, write it into metadata

```
m_al_mean <- rowMeans(expr[, -ref_cells]) # 惡性平均 profile
corr_mal <- apply(expr, 2, cor, y = m_al_mean)

s_hi <- quantile(cnv_score[ref_cells], 0.95)
lab <- ifelse(cnv_score > s_hi & corr_mal > 0.4, "malignant",
             ifelse(cnv_score <= s_hi & corr_mal < 0.2, "normal",
                    "ambiguous"))
table(lab)
```

```
lab
ambiguous malignant normal
      7      117      60
```

With a Seurat object, hang this lab vector back on with AddMetaData and carry on downstream

What this section did

Quantify



one CNV score and
one malignant-profile
correlation per cell

Cross-check



score, clustering and
markers vote together

Record



malignant / normal /
ambiguous lands in
metadata — no forced
calls

07

Where people get hurt

Three traps — each computed for real

Trap 1: the wrong reference

What it is

Putting a suspect cluster into the reference, or letting tumour cells contaminate it.



Why it happens


With shaky annotations, the biggest, most normal-looking epithelial cluster gets picked as baseline.

What it costs you

The tumour's CNV hides inside the baseline: the map zeroes out and normal cells sprout fake gains — the signs flip.


Computed: wrong baseline, flattened map

correct: immune + oligodendrocytes as reference



the real events are all visible

wrong: subclone A used as reference



malignant rows flatten to zero; normal rows show fake gains — the loss hides in the baseline and the sign flips

simulated data

Same cells, same algorithm, only the reference changed — compare the panels

Trap 2: reading a gene cluster as CNV

What it is

Co-located, co-regulated clusters like HLA or IG loci upregulate together and smooth into a fake amplification.



Why it happens


The sliding window assumes neighbouring genes vary independently — co-regulated clusters break exactly that.

What it costs you

Activated immune cells grow a fake chr6p21 event, plasma cells light up an IG locus — and it gets published as a tumour clone.

Computed: the HLA-cluster false signal

simulated data



high-risk gene clusters

HLA (chr6p21)
IG loci (chr2 / 14 / 22)
proliferation / heat-shock blocks
co-located, co-regulated, big shifts
→ smoothing turns them into fake CNV

**fix: mask these regions,
or never trust them alone**

the sliding window assumes neighbouring genes are independent — co-regulated clusters break exactly that

These 80 cells are genomically normal — mask such regions, or never trust them alone

Trap 3: clean CNV, so not cancer

What it is

Seeing a flat CNV profile and labelling the cluster normal, case closed.



Why it happens

Intuition says cancer means genomic chaos — common, yes; necessary, no.


What it costs you

Near-diploid leukaemias and some paediatric tumours get waved through as normal — downstream is wrong from step one.

CNV is evidence, not the verdict

simulated data

a spotless CNV profile... so not cancer?



chromosome 1 → 22 (near-diploid cells, computed)

near-diploid cancers are real

- most leukaemias: fusion/point-mutation driven
- some paediatric tumours: strikingly quiet genomes
- near-diploid adult tumours exist too

**CNV runs one way: events imply malignant;
'no events implies normal' does not hold**

no signal? fall back on mutations, fusions, markers, clinics

CNV is one piece of evidence, not the verdict

The inference runs one way: events imply malignant; no events never proves normal

TRAPS, SUMMED UP

**Every colour on that heatmap is an inference
relative to the baseline you chose.**

State your baseline, your window and your platform — then your CNV conclusions can stand.

5+

Strategy Room: Complex Samples

The standard route is done — this episode you already held a tumour. Now imagine it were a different one

Different tumours, different CNV fingerprints

Look up the fingerprint before you open the map — so you know what to look for, and what 'not found' means

Tumour	Textbook events	Reading decision
GBM	whole-chr7 gain + whole-chr10 loss	whole-chromosome scale — windowing's best case; highly recurrent across patients
Oligodendroglioma	1p/19q co-deletion	arm-level, this episode's example — diagnostic-grade evidence
Breast cancer	focal HER2 (ERBB2, chr17q12) amplification; luminal often 1q+ / 16q-	focal events are narrower than the 101-gene window and get diluted — leave arm-level to inferCNV, corroborate focal amps elsewhere
Liver cancer	1q+ and 8q+ (8q carries MYC)	mostly arm-level gains; the hepatocyte metabolic background is noisy — count only events that recur across patients

Near-diploid tumours: the list and the fallbacks

Whose CNV is often clean



many blood cancers
(AML can be near-normal karyotype),
some paediatric brain tumours, a few near-diploid solid tumours

Genomic fallbacks



point mutations (full-length Smart-seq reads them; 10x sees only the 3' end), fusion genes, allelic imbalance (Numbat / HoneyBADGER)


Phenotypic fallbacks



immunophenotype and the differentiation block — compare against a normal marrow / tissue atlas, look for stalled progenitors

Multi-region sampling: subclones in time and space

simulated data



shared across regions

1p- / 19q-
= early trunk events

private to one region


core B's chr4- / 11+
= late branches

space dates the evolution

margin's extra unresolved
is biology, not failure

Events shared across regions are the early trunk; region-private ones are late branches. The infiltrating margin gives weak signal — extra unresolved there is biology

Triangulation, reinforced: veto, self-check, switch



the draft is not the end — labels earn their place only after all three checks

(schematic flow; the score and cutoff mechanics are on slides 29-31)

Three insurance checks after the draft verdict: the lineage veto, the reference self-check, and switching evidence tracks when CNV is absent

Episode checklist



you can say why clustering + markers cannot call malignancy



references come from the same sample's immune / stromal cells



cutoff matches the platform: 1 for Smart-seq, 0.1 for 10x



read the heatmap reference-first, events second



verdicts triangulate CNV, clustering and markers



never trust an HLA / IG region 'event' on its own

Tick all six before a malignant label goes into your paper

ONE-LINE SUMMARY

**Malignancy calls read the genome's shadow:
order, window, baseline — then triangulate.**

CNV is strong evidence — provided you know the three places it lies.

Check yourself 1

Why can't clustering plus markers settle malignant vs normal epithelium?

- A Clustering algorithms fail on tumour data
- B Malignancy is defined genomically; both tools only see the transcriptome
- C Epithelial cells have no usable markers
- D Malignant cells are too rare

Answer B. Malignant cells share lineage markers with normal epithelium, and clusters reflect expression differences, not genomic state. You need genome-level evidence like CNV inferred from the transcriptome.

Check yourself 2

What is the sliding-window average in inferCNV for?

- A Speeding up the computation
- B Crushing single-gene noise so chromosome-level coherent shifts emerge
- C Correcting batch effects
- D Removing low-quality cells

Answer B. Per-gene ratios carry about one log2 unit of noise — no event is visible. Averaging tens of genes per window shrinks the noise several-fold, and contiguous copy-number shifts surface.

Check yourself 3

One cluster: low CNV score, mixes across patients, immune markers. Another: chr1p/19q loss, patient-specific, epithelial markers. Your labels?

- A Former normal, latter malignant
- B Both ambiguous until DNA sequencing
- C Former malignant, latter normal
- D The CNV score alone suffices

Answer A. Both are textbook three-way agreements: the first has none of the evidence (normal), the second has all three (malignant). Triangulation means calling only on agreement — and both agree.

NEXT EPISODE

B17: Deconvolution — dissecting bulk with a single-cell reference

By now you can dissect a single-cell dataset in depth. Next we flip it:
300 patients of cheap bulk data and a 3,000-cell reference — let each complete the other.